

PARAS DESHPANDE

Pune, India | +91 8788644033 | paraspdeshpande@gmail.com

EDUCATION

B.E. – Artificial Intelligence & Machine Learning

CGPA: 8.0

Savitribai Phule Pune University | June 2025 |

TECHNICAL SKILLS

Languages: Python, HTML, CSS

Libraries: NumPy, Pandas, Matplotlib, Scikit-learn, TensorFlow, PyTorch, Keras, Gensim, SpaCy, NLTK, LangChain,

AI/LLM: Prompt Engineering, RAG Pipelines, Embeddings, NLP

Tools: Streamlit, FastAPI, Docker, Linux, Excel, Power BI, langflow, n8n, make

Databases: SQL, ChromaDB, FAISS, PostgreSQL, MongoDB, Redis, Pinecone, Chroma, GCP Vector search

Other: API Integration,

Frameworks: React.js, LangChain, LangGraph, Pydantic AI, ADK (Agent Development Kit)

LLMs & Agents: OpenAI Claude, Gemini, Mistral, LLaMA

Protocols/APIs: REST, gRPC, WebSockets, OAuth2, OpenAPI, MCP, A2A

EXPERIENCE

Data Science Intern

Geta.AI Labs Pvt Ltd, Pune |

Jun 2025 – Sept 2025

- Engineered multi-LLM RAG pipelines (OpenAI, Claude, Gemini, LLaMA) boosting retrieval quality by 40%.
- Developed production-grade chatbots with LangChain, ChromaDB, PostgreSQL, Redis, MongoDB.
- Deployed modular FastAPI + Streamlit architecture with <200ms API latency and automated JSON logging.
- Integrated external APIs enabling scalable, context-aware conversational systems.

Founding Member

CoDE – Club of Data Engineers, MCOE Pune |

Oct 2023 – Nov 2024

- Led and executed 6+ technical events, including cybersecurity seminars and a 3-day Python workshop.
- Managed a 30-member team focused on LLMs, blockchain, and data engineering.
- Built assessment platform for 50+ participants; coordinated industry interactions and guest sessions.

PROJECTS

AI-Powered Business Analytics Dashboard

Python, FastAPI, Streamlit, Scikit-learn, TF-IDF, Prophet

- Built predictive models: Churn (89% accuracy), Sentiment (85%), Sales Forecasting (MAPE 7.4%).
- Delivered real-time analytics via FastAPI–Streamlit integration with sub-200ms responses.

Sales Insights Dashboard (Power BI)

Power BI, DAX, Power Query

- Identified reps with >15% variance from target.
- Built forecasting model (R^2 : 0.88) and dynamic KPIs for management decisions.

RAG Chatbot

Python, FastAPI, LangChain, Ollama, ChromaDB

- Achieved 92% retrieval accuracy with semantic chunking + embeddings.
- 55% reduction in hallucinations and 40% faster response latency.
- Built frontend supporting real-time PDF ingestion and automated knowledge base creation.

CERTIFICATIONS

- Python for Everybody – University of Michigan
- AWS Data Analytics, AI/ML
- Git & GitHub – Microsoft